Mercury: Bringing Efficiency to Key-value Stores

Rohan Gandhi, Y. Charlie Hu
Purdue University

Aayush Gupta, Anna Povzner,
Wendy Belluomini, Tim Kaldewey
IBM Research – Almaden

Overview:
In-memory Key-Value Stores

- Store data in DRAM
- Alleviate database load

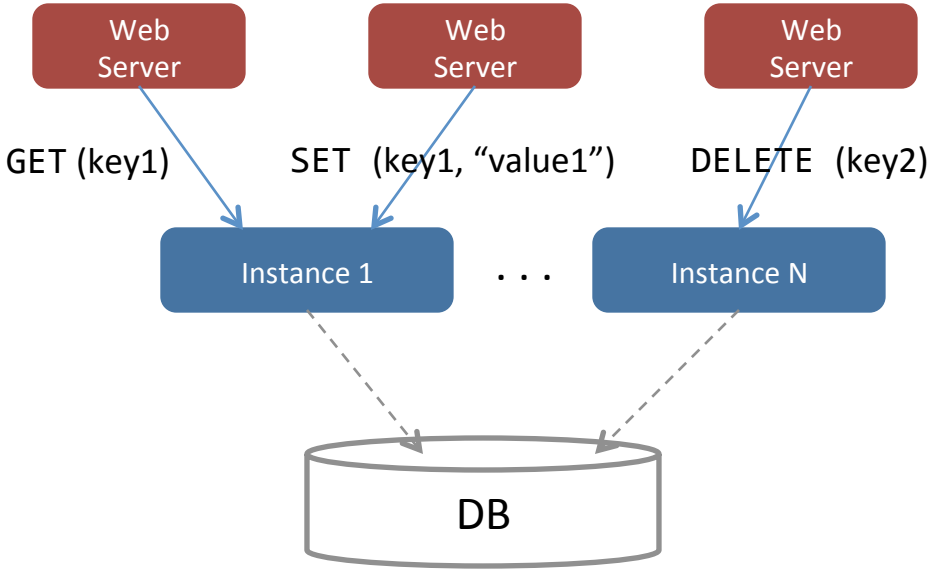

Memcached	MemC3	MassTree
http://memcached.org/	NSDI' 13	EuroSys' 12

Request Latency = Latency(Network) + Latency(Server)

Bottleneck

Memory Efficiency Matters!

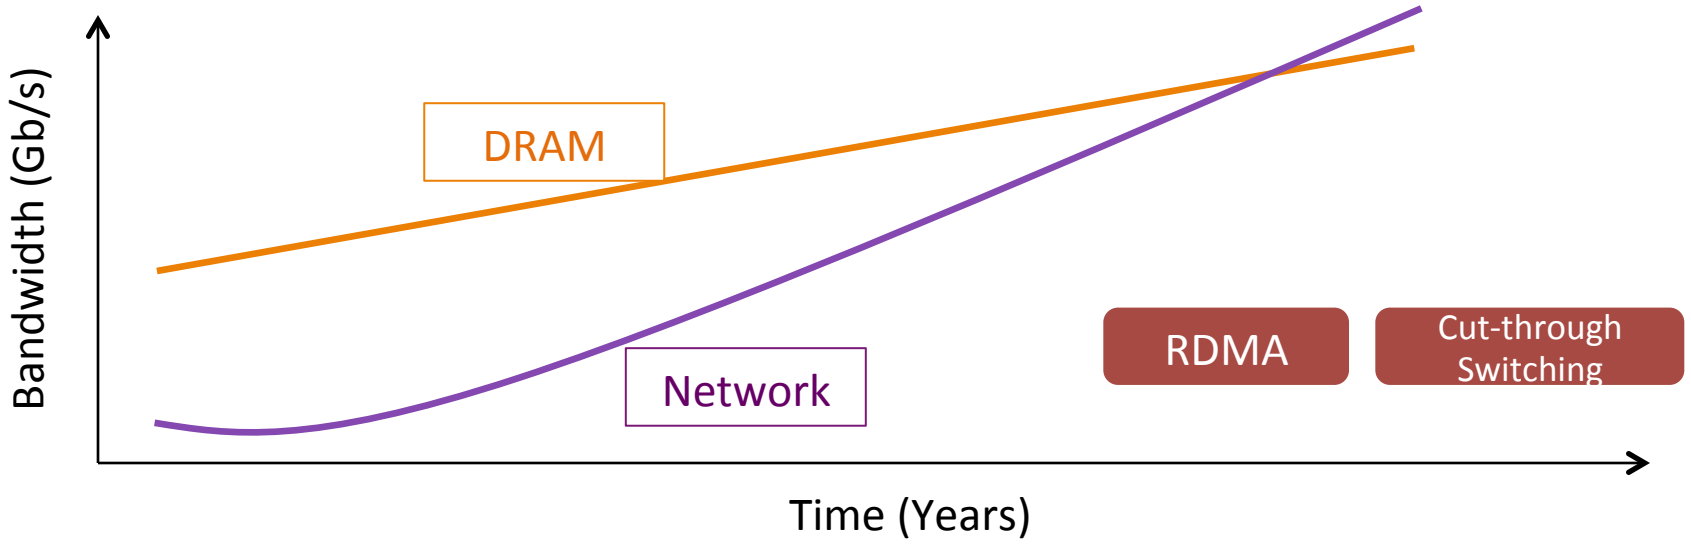

Bottlenecks shifting from network to DRAM !

Request Latency = Latency(Network) + Latency(Server)

Bottleneck

But we have CPU caches (L1:L3) to alleviate DRAM performance !

Is it true with key-value stores ?

Workloads:
Key-values pairs are becoming small

Facebook Trace: Sigmetrics'12

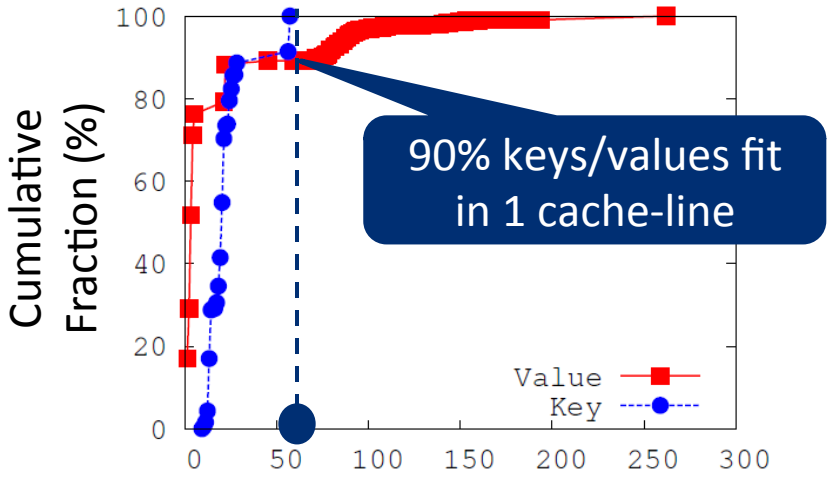

Write-heavy

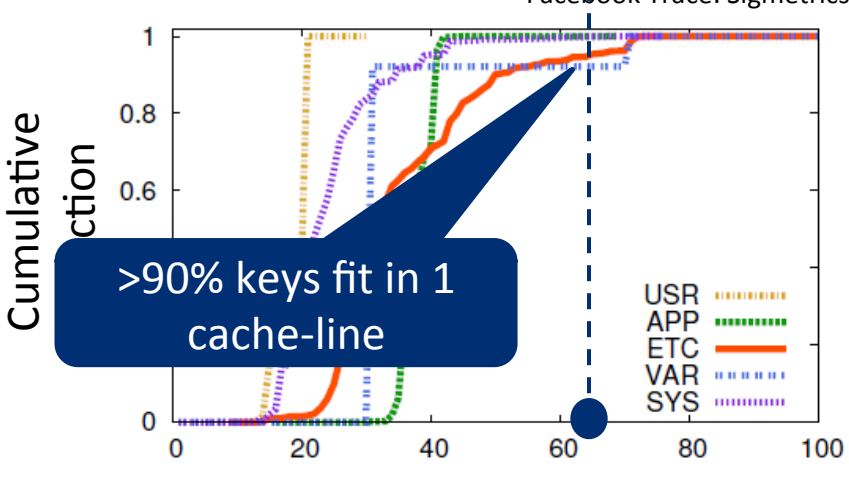

Read-heavy

Implications:

- No cache prefetching
- Working set > cache size
- Large number of key-value pairs / node
- Metadata > cache size

Multi-threading → Synchronization overheads

Performance Comparison

	~Zero Synchronization Overhead ?	Minimum DRAM Accesses ?	Workload Independent ?
Memcached	**No**	Yes	Yes
MassTree (Eurosys'12)	Yes	**No**	Yes
MemC3 (NSDI'13)	Yes	Yes	**No**

Mercury	**Yes**	**Yes**	**Yes**

Mercury

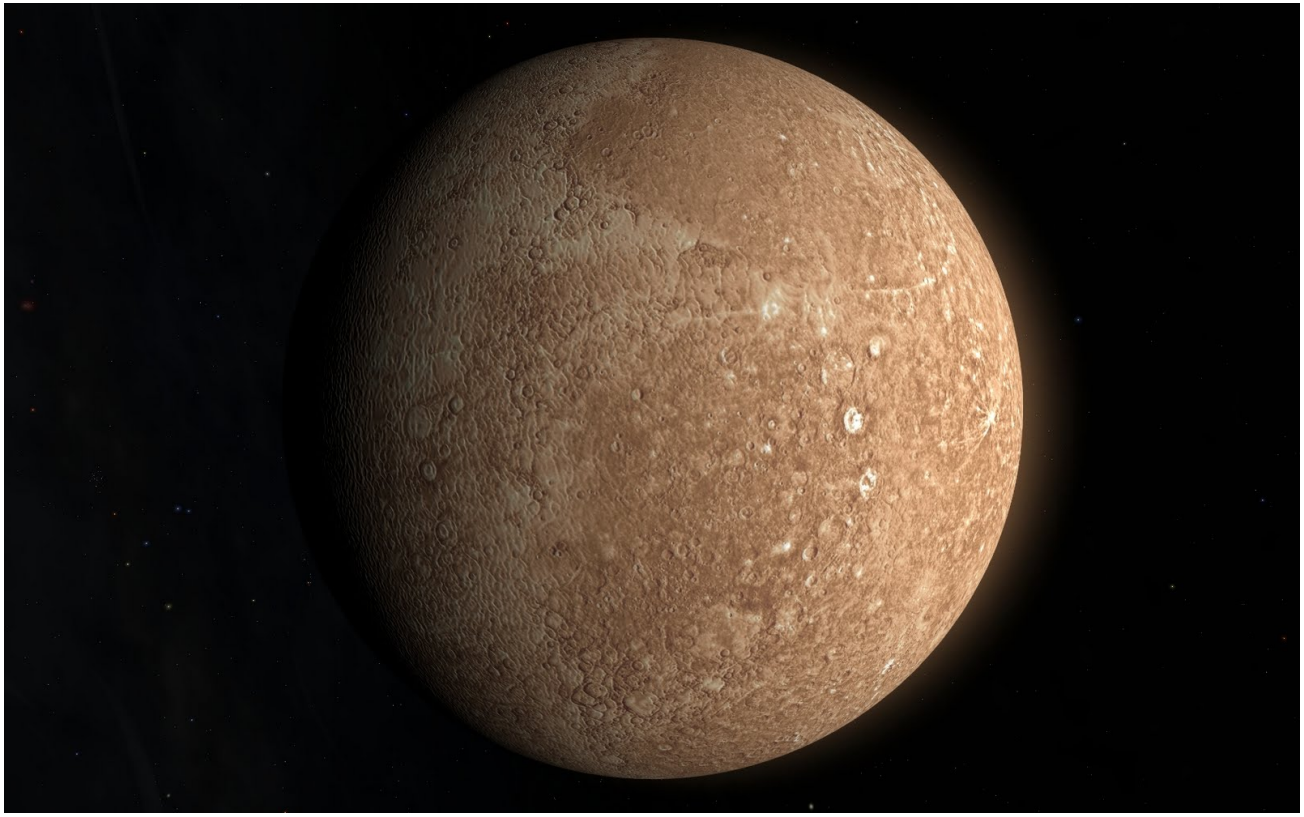

Closest planet to sun !

Mercury:
Low latency-High Throughput Key-Value Store

- Improves Memory system for a single server
- Operations: GET/SET/DELETE

- Workload independent
- Scalability
  - Store millions of key-value pairs
  - Support multiple threads and multiple writers
- Minimum DRAM accesses
  - 1 access: Hash-table
  - 1 access: Key-value pair

Design Choices

Lookup: Hash-table or Tree?

Chain Hashing

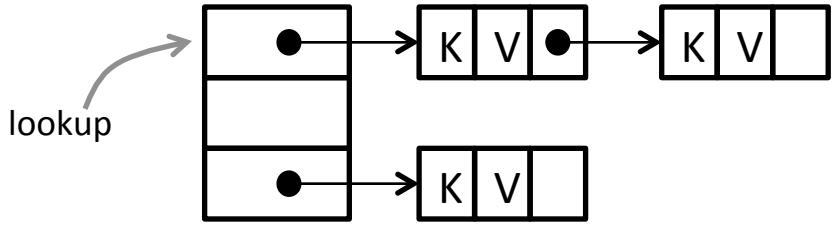

Advantage:
Lock granularity: Single bucket

Challenge:
Keep chain length small

Linear Probing

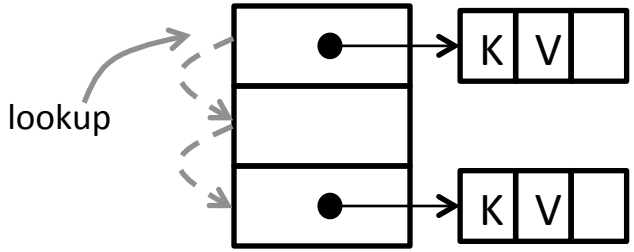

Advantage:
Cache-friendly

Challenge:
Locking ?

Cuckoo Hashing

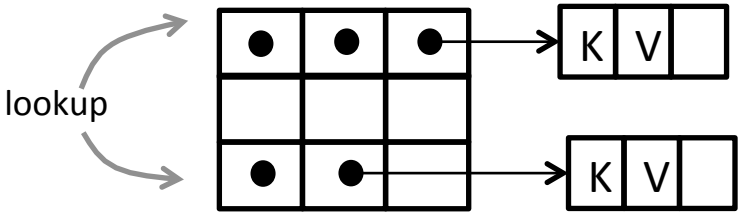

Advantage:
Higher hash-table occupancy

Challenge:
Locking? Expansion?

Mercury uses chain hashing
Dynamically expands hash-table size to keep chain lengths small

Mercury: Core data-structure

Chained Hash-table + each bucket protected with a single lock

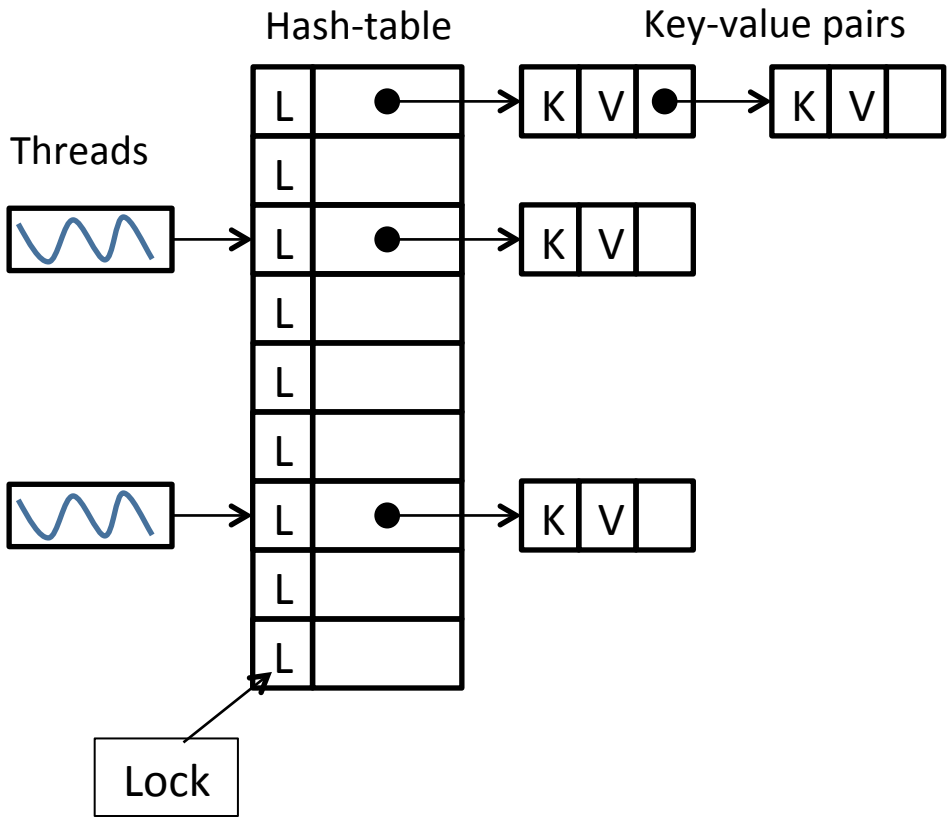

Operations:
1) Acquire lock
2) Read/update hash-table entry

1 DRAM access due to 64-byte cache line

Contention probability@
(1 million locks, 12 thread)
= 0.07×10^{-3}

All locks implemented using TAS
(not pthread)

How to expand hash-table?
How to lookup key-value pairs during expansion?

Design: Hash-table Expansion

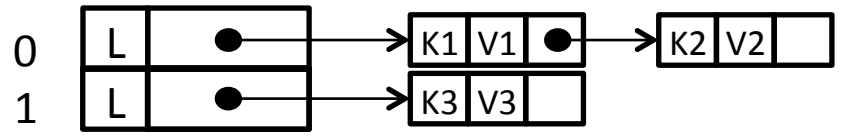

Expand-thread

Moves key-value pairs to bigger hash-table

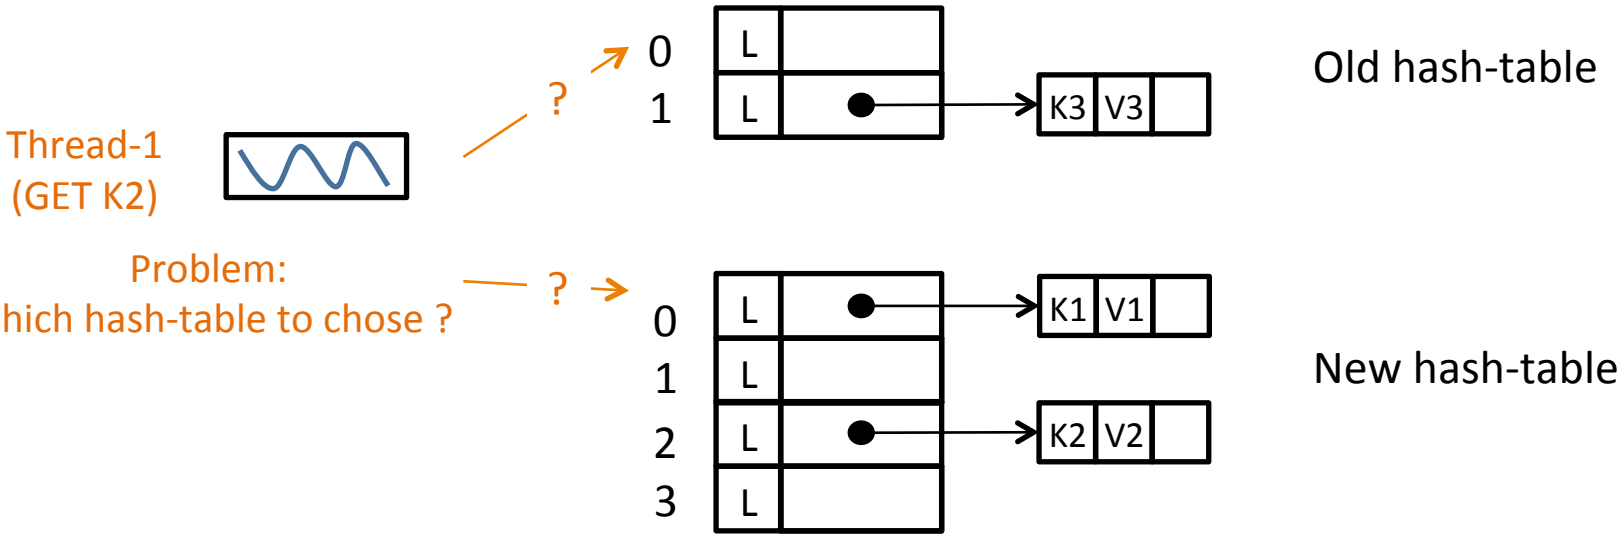

How to achieve lookup during hash-table expansion ?

Design: Lookup During Expansion

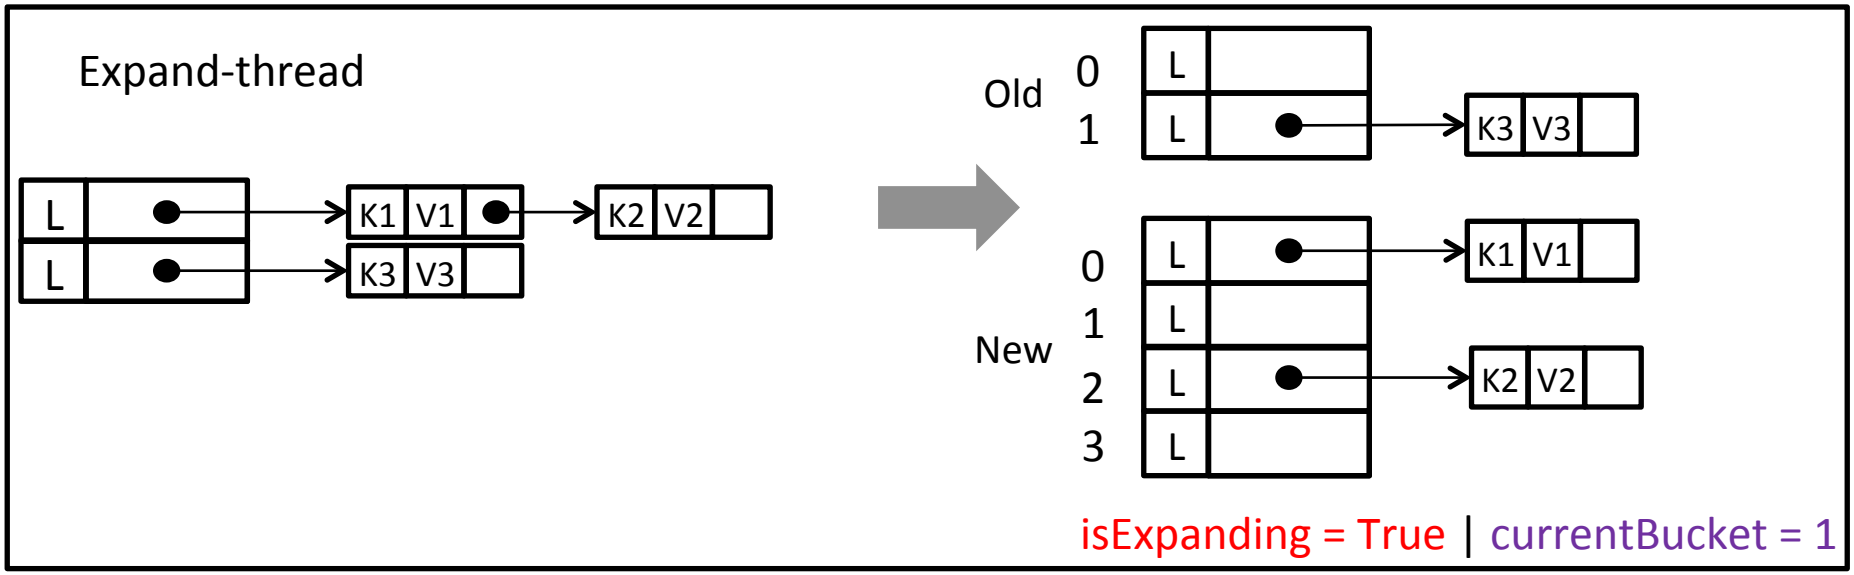

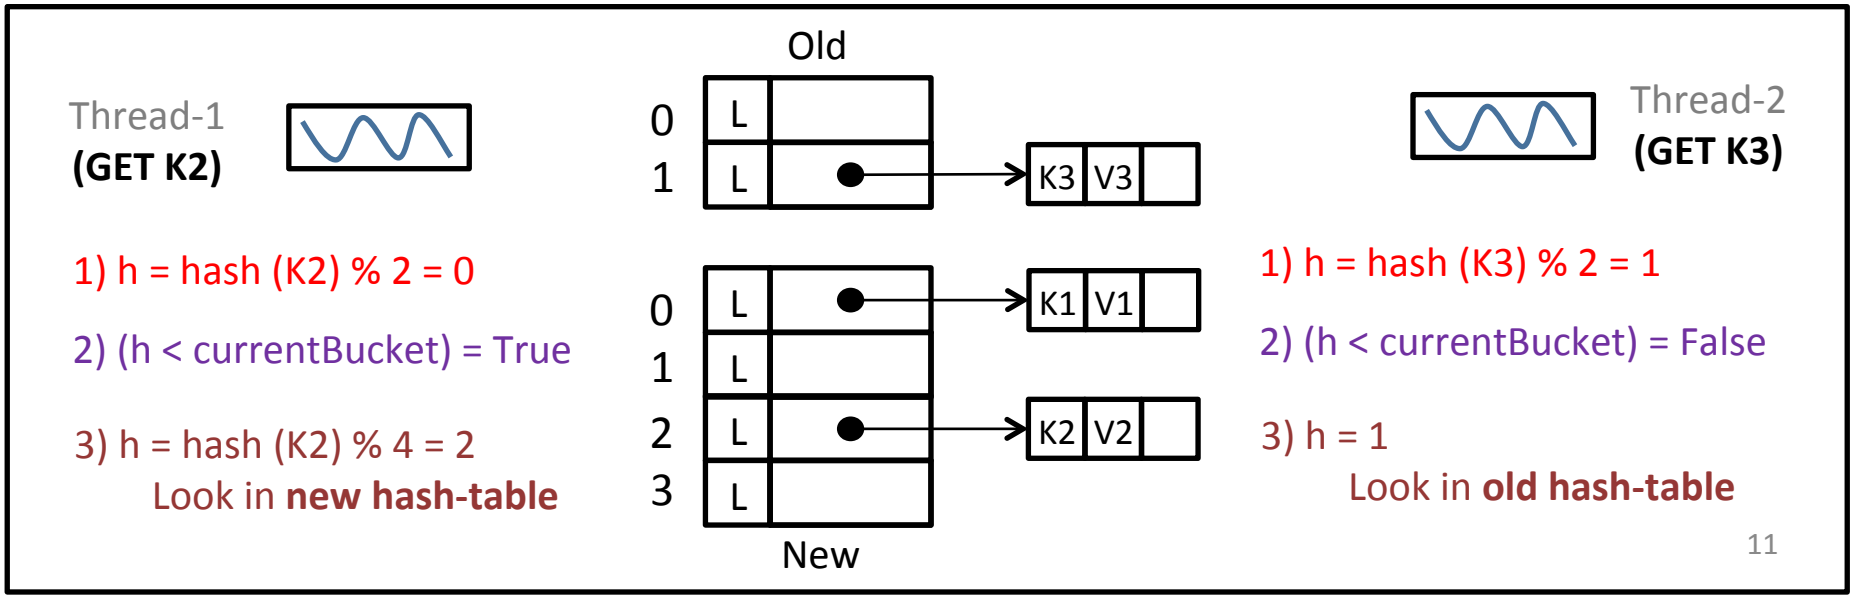

Expanding Hash-table: Race Condition

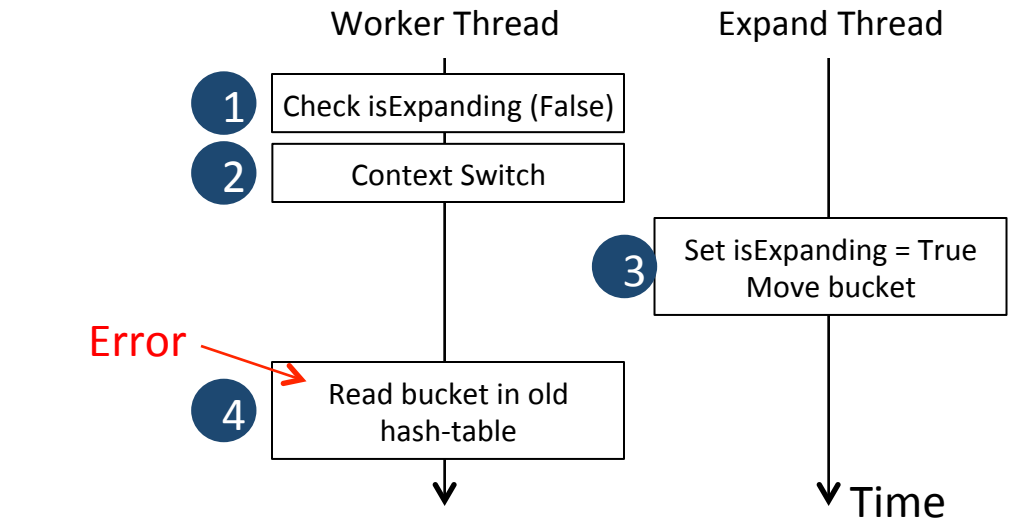

isExpanding bit needs synchronization

Contention: <1 usec

No cache invalidations between threads

hash-table size Independent

Design: Memory Management

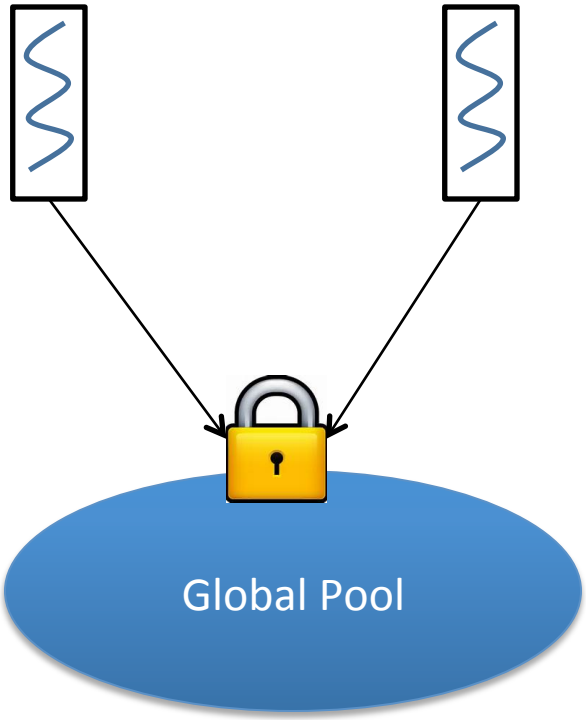

Lock Contention

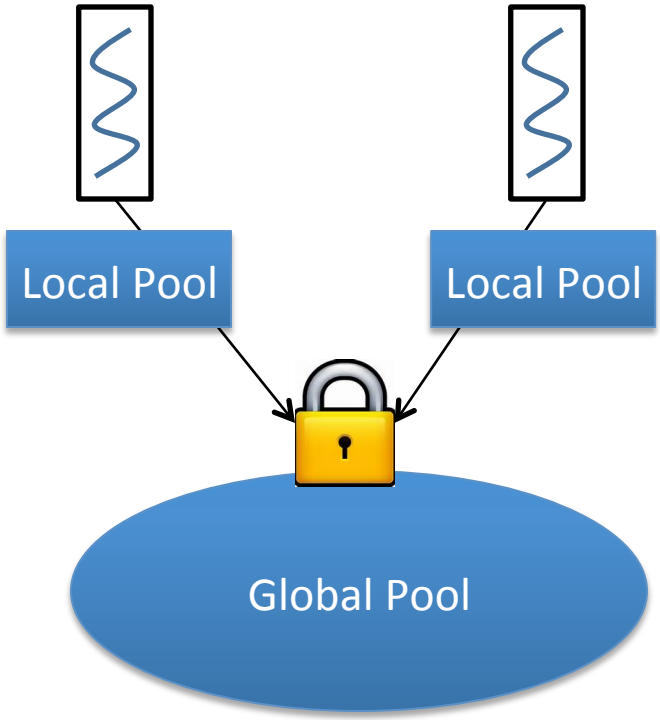

Lock Contention Reduced

Mercury Design Summary

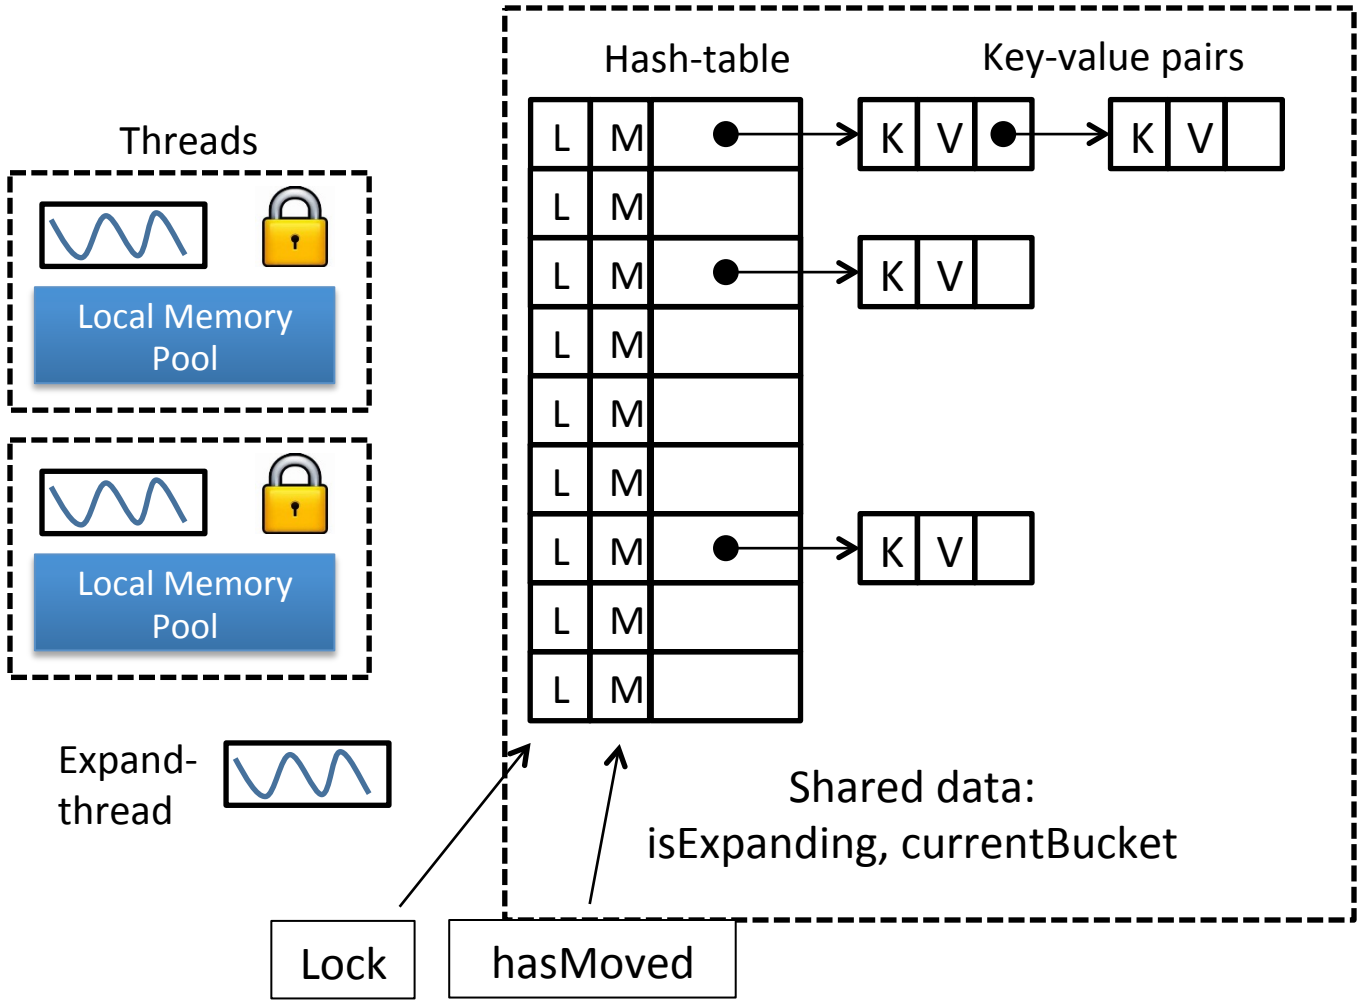

Evaluation: Cache Misses

Setup:
Key/Value size: 8-bytes each
Dataset: 32 million key-value pairs
#Queries: 100 million

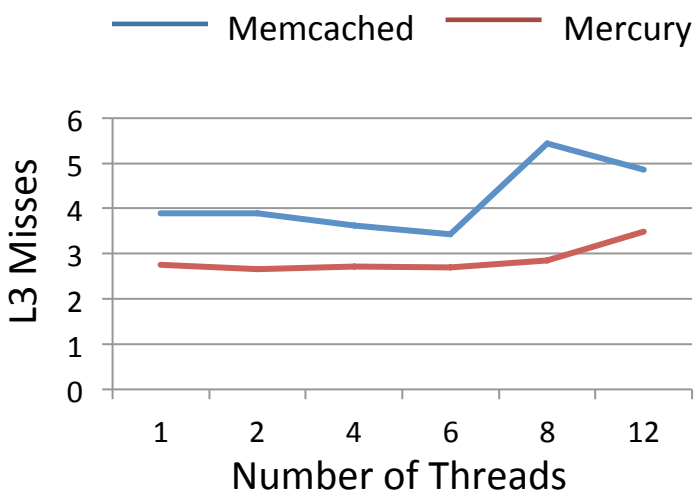

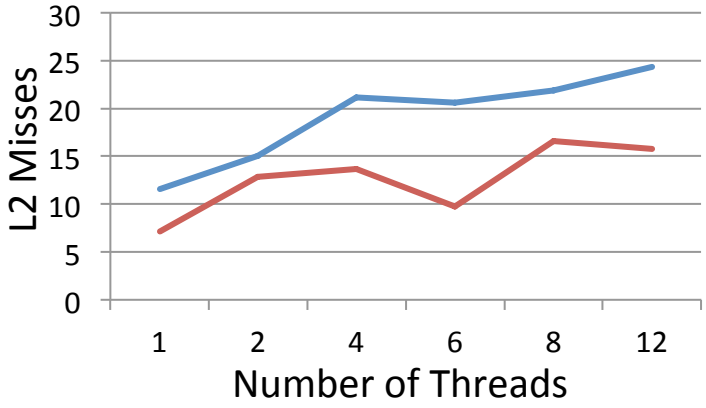

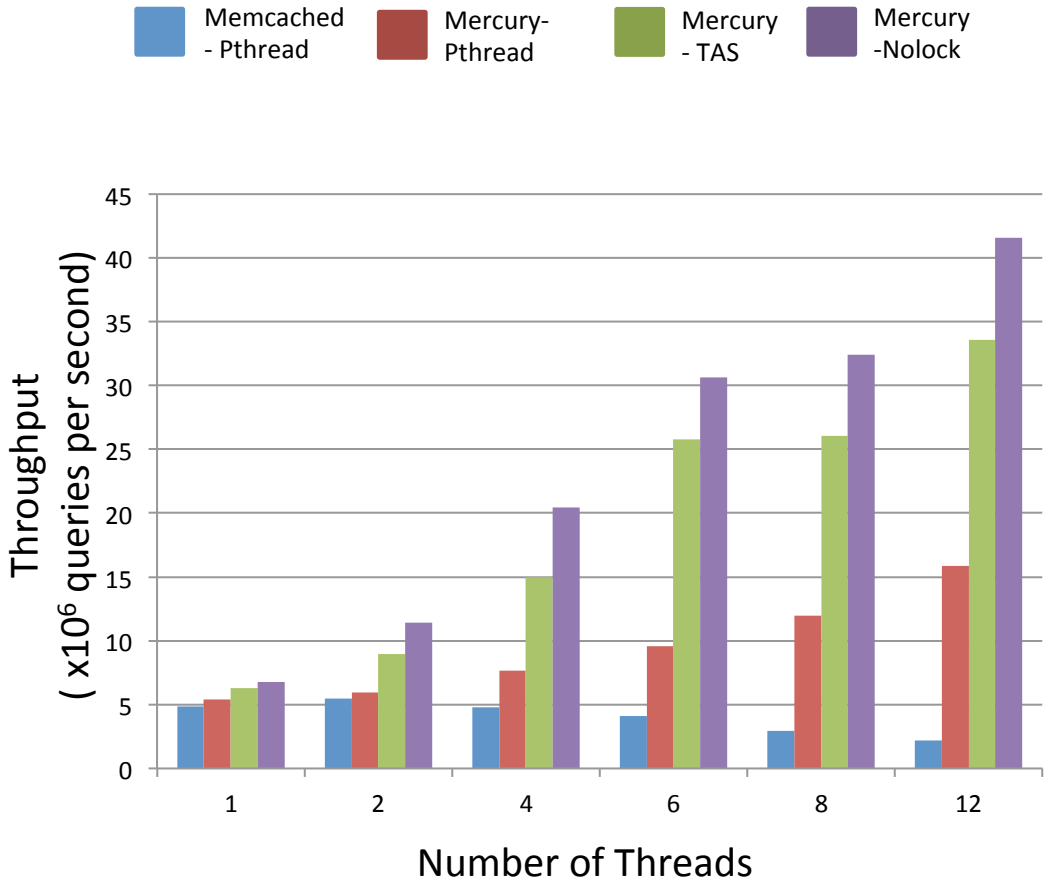

Evaluation: YCSB

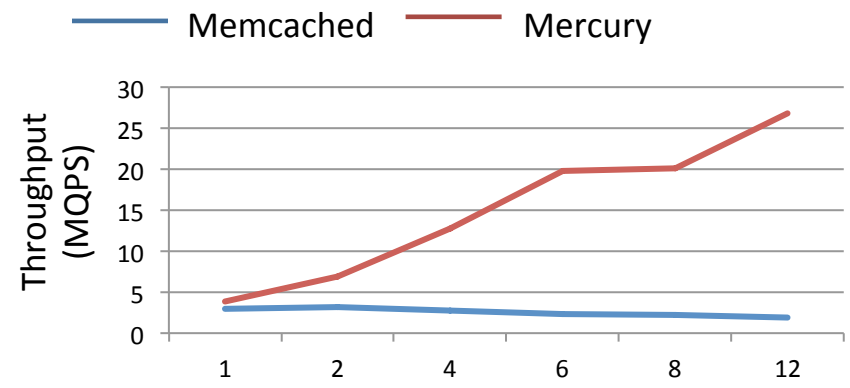

100% GET

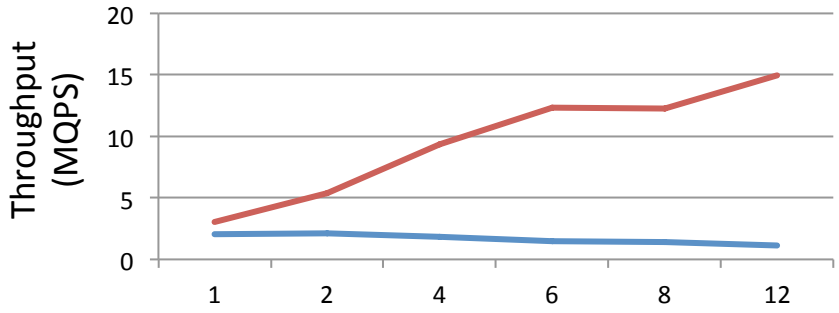

75% GET

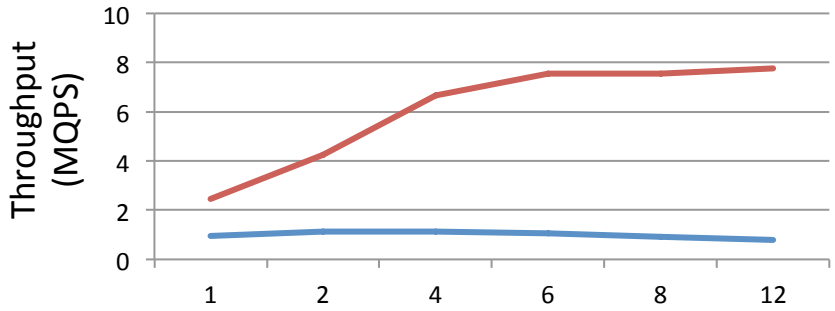

50% GET

YCSB Setup:

Key size: 24-bytes
Value size: 10-byte

Dataset: 24 million key-values
#Queries: 100 million

Conclusion

- Mercury
  - Memory efficient key-value store
  - Improves Memcached throughput by 14x
  - Workload Independent
  - Scales for number of key-value pairs and threads

- Design Choices
  - Chain-hashing (with expansion): 2 DRAM Access latency
  - Fine-grained locking: ~Zero contention

Thank You!

Special thanks to Pin Zhou, Ronen Kat, Sivan Toledo, Mary Baker and remaining organization committee !